

Exercise 8

Fundamentals of machine vision algorithms

dr.sc. Filip Šuligoj
fsuligoj@fsb.hr




University of
Zagreb




Faculty of mechanical
engineering and naval
architecture

Practical task 1

Template matching



Test different methods and templates



Practical task 1

Template matching

◆ matchTemplate()

```
void cv::matchTemplate ( InputArray image,
                        InputArray templ,
                        OutputArray result,
                        int method,
                        InputArray mask = noArray()
                      )
```

Python:

```
cv.matchTemplate( image, templ, method[, result[, mask]] ) -> result
```

```
#include <opencv2/imgproc.hpp>
```

Compares a template against overlapped image regions.

The function slides through *image*, compares the overlapped patches of size $w \times h$ against *templ* using the specified method and stores the comparison results in result. Here are the formulae for the available comparison methods (*I* denotes image, *T* template, *R* result). The summation is done over template and/or the image patch: $x' = 0 \dots w - 1, y' = 0 \dots h - 1$

After the function finishes the comparison, the best matches can be found as global minimums (when **TM_SQDIFF** was used) or maximums (when **TM_CCORR** or **TM_CCOEFF** was used) using the **minMaxLoc** function. In case of a color image, template summation in the numerator and each sum in the denominator is done over all of the channels and separate mean values are used for each channel. That is, the function can take a color template and a color image. The result will still be a single-channel image, which is easier to analyze.

Parameters

image Image where the search is running. It must be 8-bit or 32-bit floating-point.
templ Searched template. It must be not greater than the source image and have the same data type.
result Map of comparison results. It must be single-channel 32-bit floating-point. If image is $W \times H$ and *templ* is $w \times h$, then result is $(W - w + 1) \times (H - h + 1)$.
method Parameter specifying the comparison method, see [TemplateMatchModes](#)
mask Mask of searched template. It must have the same datatype and size with *templ*. It is not set by default. Currently, only the **TM_SQDIFF** and **TM_CCORR_NORMED** methods are supported.

Examples:

```
samples/cpp/tutorial_code/Histograms_Matching/MatchTemplate_Demo.cpp.
```

◆ normalize()

```
void cv::normalize ( InputArray src,
                    InputOutputArray dst,
                    double alpha = 1,
                    double beta = 0,
                    int norm_type = NORM_L2,
                    int dtype = -1,
                    InputArray mask = noArray()
                  )
```

Python:

```
cv.normalize( src, dst[, alpha[, beta[, norm_type[, dtype[, mask]]]] ) -> dst
```

```
#include <opencv2/core.hpp>
```

Normalizes the norm or value range of an array.

The function `cv::normalize` normalizes scale and shift the input array elements so that

$$\|dst\|_{p_r} = \alpha$$

(where $p = \text{inf}, 1$ or 2) when `normType=NORM_INF, NORM_L1, or NORM_L2`, respectively; or so that

$$\min dst(I) = \alpha, \max dst(I) = \beta$$

◆ minMaxLoc()

```
void cv::minMaxLoc ( InputArray src,
                    double * minVal,
                    double * maxVal = 0,
                    Point * minLoc = 0,
                    Point * maxLoc = 0,
                    InputArray mask = noArray()
                  )
```

Python:

```
cv.minMaxLoc( src[, mask] ) -> minVal, maxVal, minLoc, maxLoc
```

```
#include <opencv2/core.hpp>
```

Finds the global minimum and maximum in an array.

The function `cv::minMaxLoc` finds the minimum and maximum element values and their positions. The extremums are searched across the whole array or, if *mask* is not an empty array, in the specified array region.

The function do not work with multi-channel arrays. If you need to find minimum or maximum elements across all the channels, use [Mat::reshape](#) first to reinterpret the array as single-channel. Or you may extract the particular channel using either `extractImageCOI`, or `mixChannels`, or `split`.

Parameters

src input single-channel array.
minVal pointer to the returned minimum value; NULL is used if not required.
maxVal pointer to the returned maximum value; NULL is used if not required.
minLoc pointer to the returned minimum location (in 2D case); NULL is used if not required.
maxLoc pointer to the returned maximum location (in 2D case); NULL is used if not required.
mask optional mask used to select a sub-array.

See also

[max](#), [min](#), [compare](#), [inRange](#), [extractImageCOI](#), [mixChannels](#), [split](#), [Mat::reshape](#)

Examples:

```
samples/cpp/image_alignment.cpp, samples/cpp/tutorial_code/Histograms_Matching/MatchTemplate_Demo.cpp,
samples/dnn/classification.cpp, samples/dnn/object_detection.cpp, samples/dnn/openpose.cpp, and samples/dnn/text_detection.cpp.
```



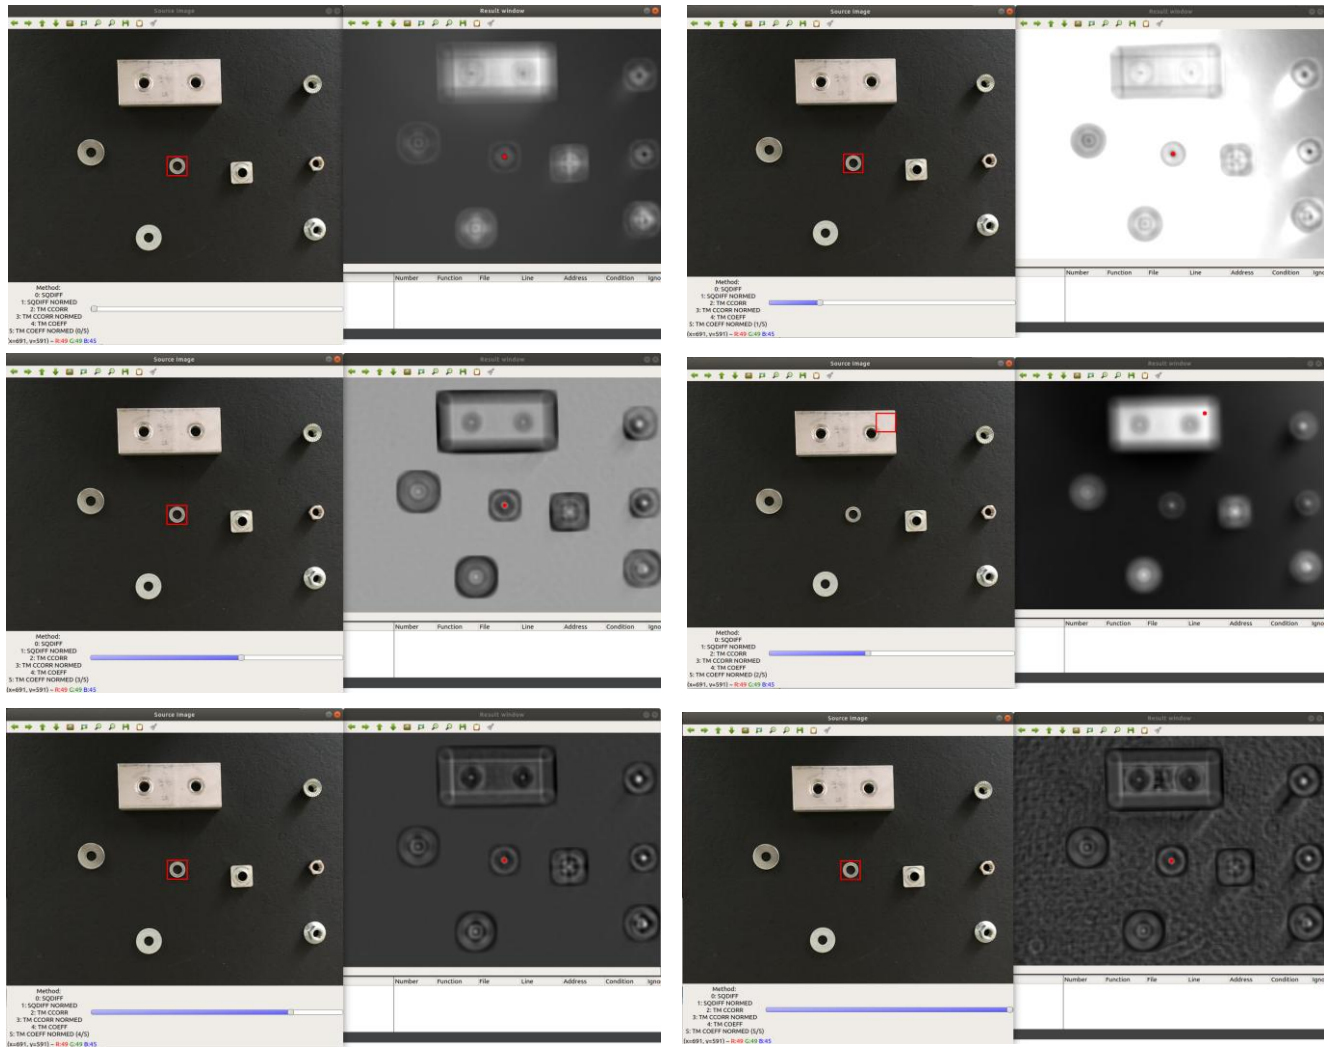
University of
Zagreb





Faculty of mechanical
engineering and naval
architecture

Practical task 1

Template matching - example of results



Practical task 2



Practical task 2

Purpose: Converts a rotation vector into a 3×3 rotation matrix.

◆ Rodrigues()

```
void cv::Rodrigues ( InputArray  src,
                    OutputArray dst,
                    OutputArray jacobian = noArray()
                  )
```

Python:

```
cv.Rodrigues( src[, dst[, jacobian]] ) -> dst, jacobian
```

```
#include <opencv2/calib3d.hpp>
```

Converts a rotation matrix to a rotation vector or vice versa.

Parameters

src Input rotation vector (3x1 or 1x3) or rotation matrix (3x3).

dst Output rotation matrix (3x3) or rotation vector (3x1 or 1x3), respectively.

jacobian Optional output Jacobian matrix, 3x9 or 9x3, which is a matrix of partial derivatives of the output array components with respect to the input array components.

$$\theta \leftarrow \text{norm}(r)$$

$$r \leftarrow r/\theta$$

$$R = \cos(\theta)I + (1 - \cos(\theta))rr^T + \sin(\theta) \begin{bmatrix} 0 & -r_z & r_y \\ r_z & 0 & -r_x \\ -r_y & r_x & 0 \end{bmatrix}$$

Inverse transformation can be also done easily, since

$$\sin(\theta) \begin{bmatrix} 0 & -r_z & r_y \\ r_z & 0 & -r_x \\ -r_y & r_x & 0 \end{bmatrix} = \frac{R - R^T}{2}$$

A rotation vector is a convenient and most compact representation of a rotation matrix (since any rotation matrix has just 3 degrees of freedom). The representation is used in the global 3D geometry optimization procedures like `calibrateCamera`, `stereoCalibrate`, or `solvePnP`.

Purpose: Applies a perspective (homography) warp to simulate a 3D transformation in 2D space.

◆ warpPerspective()

```
void cv::warpPerspective ( InputArray  src,
                           OutputArray dst,
                           InputArray  M,
                           Size        dsize,
                           int         flags = INTER_LINEAR,
                           int         borderMode = BORDER_CONSTANT,
                           const Scalar & borderValue = Scalar()
                         )
```

Python:

```
cv.warpPerspective( src, M, dsize[, dst[, flags[, borderMode[, borderValue]]]] ) -> dst
```

```
#include <opencv2/imgproc.hpp>
```

Applies a perspective transformation to an image.

The function `warpPerspective` transforms the source image using the specified matrix:

$$\text{dst}(x, y) = \text{src} \left(\frac{M_{11}x + M_{12}y + M_{13}}{M_{31}x + M_{32}y + M_{33}}, \frac{M_{21}x + M_{22}y + M_{23}}{M_{31}x + M_{32}y + M_{33}} \right)$$

when the flag `WARP_INVERSE_MAP` is set. Otherwise, the transformation is first inverted with `invert` and then put in the formula above instead of `M`. The function cannot operate in-place.

Parameters

src input image.

dst output image that has the size `dsize` and the same type as `src`.

M 3×3 transformation matrix.

dsize size of the output image.

flags combination of interpolation methods (`INTER_LINEAR` or `INTER_NEAREST`) and the optional flag `WARP_INVERSE_MAP`, that sets `M` as the inverse transformation (`dst` \rightarrow `src`).

borderMode pixel extrapolation method (`BORDER_CONSTANT` or `BORDER_REPLICATE`).

borderValue value used in case of a constant border; by default, it equals 0.

See also

`warpAffine`, `resize`, `remap`, `getRectSubPix`, `perspectiveTransform`



University of
Zagreb



Faculty of mechanical
engineering and naval
architecture

Practical task 2

Purpose: Estimates the 3D pose (rotation and translation) of an object based on known geometry and its 2D image.

◆ solvePnP()

```
bool cv::solvePnP ( InputArray  objectPoints,
                   InputArray  imagePoints,
                   InputArray  cameraMatrix,
                   InputArray  distCoeffs,
                   OutputArray  rvec,
                   OutputArray  tvec,
                   bool         useExtrinsicGuess = false ,
                   int          flags = SOLVEPNP_ITERATIVE
                 )
```

Python:

```
cv.solvePnP( objectPoints, imagePoints, cameraMatrix, distCoeffs[, rvec[, tvec[, useExtrinsicGuess[, flags]]]] ) -> retval, rvec, tvec
```

```
#include <opencv2/calib3d.hpp>
```

Finds an object pose from 3D-2D point correspondences.

Purpose: Projects 3D object points to 2D image coordinates using estimated pose.

◆ projectPoints()

```
void cv::projectPoints ( InputArray  objectPoints,
                        InputArray  rvec,
                        InputArray  tvec,
                        InputArray  cameraMatrix,
                        InputArray  distCoeffs,
                        OutputArray  imagePoints,
                        OutputArray  jacobian = noArray() ,
                        double       aspectRatio = 0
                      )
```

Python:

```
cv.projectPoints( objectPoints, rvec, tvec, cameraMatrix, distCoeffs[, imagePoints[, jacobian[, aspectRatio]]] ) -> imagePoints, jacobian
```

```
#include <opencv2/calib3d.hpp>
```

Projects 3D points to an image plane.

Parameters




- objectPoints** Array of object points expressed wrt. the world coordinate frame. A $3 \times N \times N \times 3$ 1-channel or $1 \times N \times N \times 1$ 3-channel (or vector<Point3f>), where N is the number of points in the view.
- rvec** The rotation vector (**Rodrigues**) that, together with tvec, performs a change of basis from world to camera coordinate system, see [calibrateCamera](#) for details.
- tvec** The translation vector, see parameter description above.
- cameraMatrix** Camera intrinsic matrix $A = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$.
- distCoeffs** Input vector of distortion coefficients $(k_1, k_2, p_1, p_2[, k_3[, k_4, k_5, k_6[, s_1, s_2, s_3, s_4[, \tau_x, \tau_y]]]])$ of 4, 5, 8, 12 or 14 elements . If the vector is empty, the zero distortion coefficients are assumed.
- imagePoints** Output array of image points, $1 \times N \times N \times 1$ 2-channel, or vector<Point2f> .
- jacobian** Optional output $2N \times (10 + \text{numDistCoeffs})$ jacobian matrix of derivatives of image points with respect to components of the rotation vector, translation vector, focal lengths, coordinates of the principal point and the distortion coefficients. In the old interface different components of the jacobian are returned via different output parameters.
- aspectRatio** Optional "fixed aspect ratio" parameter. If the parameter is not 0, the function assumes that the aspect ratio (f_x/f_y) is fixed and correspondingly adjusts the jacobian matrix.

The function computes the 2D projections of 3D points to the image plane, given intrinsic and extrinsic camera parameters. Optionally, the function computes Jacobians -matrices of partial derivatives of image points coordinates (as functions of all the input parameters) with respect to the particular parameters, intrinsic and/or extrinsic. The Jacobians are used during the global optimization in [calibrateCamera](#), [solvePnP](#), and [stereoCalibrate](#). The function itself can also be used to compute a re-projection error, given the current intrinsic and extrinsic parameters.



Practical task 3

SURF for feature matching – prepared code



Practical task 3

◆ create()

```
static Ptr<SURF> cv::xfeatures2d::SURF::create ( double hessianThreshold = 100 ,
                                              int nOctaves = 4 ,
                                              int nOctaveLayers = 3 ,
                                              bool extended = false ,
                                              bool upright = false
                                              )
```

Python:

```
cv.xfeatures2d.SURF_create( [ , hessianThreshold[, nOctaves[, nOctaveLayers[, extended[, upright]]]] ) -> retval
```

Parameters

hessianThreshold Threshold for hessian keypoint detector used in SURF.

nOctaves Number of pyramid octaves the keypoint detector will use.

nOctaveLayers Number of octave layers within each octave.

extended Extended descriptor flag (true - use extended 128-element descriptors; false - use 64-element descriptors).

upright Up-right or rotated features flag (true - do not compute orientation of features; false - compute orientation).

◆ knnMatch() [1/2]

```
void cv::DescriptorMatcher::knnMatch ( InputArray queryDescriptors,
                                     InputArray trainDescriptors,
                                     std::vector< std::vector< DMatch > > & matches,
                                     int k,
                                     InputArray mask = noArray() ,
                                     bool compactResult = false
                                     ) const
```

Python:

```
cv.DescriptorMatcher.knnMatch( queryDescriptors, trainDescriptors, k[, mask[, compactResult]] ) -> matches
cv.DescriptorMatcher.knnMatch( queryDescriptors, k[, masks[, compactResult]] ) -> matches
```

Finds the k best matches for each descriptor from a query set.

Parameters

queryDescriptors Query set of descriptors.

trainDescriptors Train set of descriptors. This set is not added to the train descriptors collection stored in the class object.

mask Mask specifying permissible matches between an input query and train matrices of descriptors.

matches Matches. Each matches[i] is k or less matches for the same query descriptor.

k Count of best matches found per each query descriptor or less if a query descriptor has less than k possible matches in total.

compactResult Parameter used when the mask (or masks) is not empty. If compactResult is false, the matches vector has the same size as queryDescriptors rows. If compactResult is true, the matches vector does not contain matches for fully masked-out query descriptors.

These extended variants of `DescriptorMatcher::match` methods find several best matches for each query descriptor. The matches are returned in the distance increasing order. See `DescriptorMatcher::match` for the details about query and train descriptors.

Purpose: Detects keypoints and computes descriptors in a single call (used with feature detectors like ORB, SIFT, etc.).

◆ detectAndCompute()

```
virtual void cv::Feature2D::detectAndCompute ( InputArray image,
                                              InputArray mask,
                                              std::vector< KeyPoint > & keypoints,
                                              OutputArray descriptors,
                                              bool useProvidedKeypoints = false
                                              )
```

Python:

```
cv.Feature2D.detectAndCompute( image, mask[, descriptors[, useProvidedKeypoints]] ) -> keypoints, descriptors
```

Detects keypoints and computes the descriptors

Purpose: Finds the k-nearest best matches between descriptor sets using distance metrics (usually for feature matching).



Practical task 3

Purpose: Applies a homography matrix to a set of 2D points to map them into a new perspective.

◆ perspectiveTransform()

```
void cv::perspectiveTransform ( InputArray  src,
                               OutputArray dst,
                               InputArray  m
                             )
```

Python:

```
cv.perspectiveTransform( src, m[, dst] ) -> dst
```

```
#include <opencv2/core.hpp>
```

Performs the perspective matrix transformation of vectors.

The function `cv::perspectiveTransform` transforms every element of `src` by treating it as a 2D or 3D vector, in the following way:

$$(x, y, z) \rightarrow (x'/w, y'/w, z'/w)$$

where

$$(x', y', z', w') = \text{mat} \cdot \begin{bmatrix} x & y & z & 1 \end{bmatrix}$$

and

$$w = \begin{cases} w' & \text{if } w' \neq 0 \\ \infty & \text{otherwise} \end{cases}$$

Here a 3D vector transformation is shown. In case of a 2D vector transformation, the z component is omitted.

Note

The function transforms a sparse set of 2D or 3D vectors. If you want to transform an image using perspective transformation, use `warpPerspective`. If you have an inverse problem, that is, you want to compute the most probable perspective transformation out of several pairs of corresponding points, you can use `getPerspectiveTransform` or `findHomography`.

Parameters

src input two-channel or three-channel floating-point array; each element is a 2D/3D vector to be transformed.

dst output array of the same size and type as `src`.

m 3x3 or 4x4 floating-point transformation matrix.

See also

`transform`, `warpPerspective`, `getPerspectiveTransform`, `findHomography`

Purpose: Computes the best-fit 3x3 homography matrix that maps one set of 2D points to another.

◆ findHomography() [1/2]

```
Mat cv::findHomography ( InputArray  srcPoints,
                        InputArray  dstPoints,
                        int          method = 0,
                        double       ransacReprojThreshold = 3,
                        OutputArray  mask = noArray(),
                        const int     maxiters = 2000,
                        const double  confidence = 0.995
                      )
```

Python:

```
cv.findHomography( srcPoints[, dstPoints[, method[, ransacReprojThreshold[, mask[, maxiters[, confidence]]]]]) -> retval, mask
```

```
#include <opencv2/calib3d.hpp>
```

Finds a perspective transformation between two planes.

Parameters

srcPoints Coordinates of the points in the original plane, a matrix of the type CV_32FC2 or vector<Point2f>.
dstPoints Coordinates of the points in the target plane, a matrix of the type CV_32FC2 or a vector<Point2f>.
method Method used to compute a homography matrix. The following methods are possible:

- 0 - a regular method using all the points, i.e., the least squares method
- RANSAC - RANSAC-based robust method
- LMEDS - Least-Median robust method
- RHO - PROSAC-based robust method

ransacReprojThreshold Maximum allowed reprojection error to treat a point pair as an inlier (used in the RANSAC and RHO methods only). That is, if

$$\| \text{dstPoints}_i - \text{convertPointsHomogeneous}(H * \text{srcPoints}_i) \|_2 > \text{ransacReprojThreshold}$$

then the point i is considered as an outlier. If `srcPoints` and `dstPoints` are measured in pixels, it usually makes sense to set this parameter somewhere in the range of 1 to 10.

mask Optional output mask set by a robust method (RANSAC or LMEDS). Note that the input mask values are ignored.

maxiters The maximum number of RANSAC iterations.

confidence Confidence level, between 0 and 1.



Student assignment (seminar)



University of
Zagreb



Faculty of mechanical
engineering and naval
architecture

Student assignment - seminar

- *Find ArUco Markers – Detect at least four ArUco markers in an input image (e.g., photo of a printed A4 sheet with markers at the corners).*
- *Rectify Image – Use the markers to compute a homography and warp the perspective, rectifying the image to a top-down view.*
- *Crop Image – Automatically crop the rectified image based on the marker layout without knowing the original aspect ratio.*
- *Threshold Image – Convert the rectified image to grayscale and apply adaptive binary thresholding to highlight darker objects.*
- *Show Contours – Detect and draw contours of all dark shapes on a white background.*

