

Exercise – OpenCV basics

Machine vision algorithms

Doc.dr.sc. Filip Šuligoj
fsuligoj@fsb.hr



University of
Zagreb



Faculty of mechanical
engineering and naval
architecture

OpenCV modules

OpenCV Overview

OpenCV (Open Source Computer Vision Library) is a widely used open-source library offering hundreds of computer vision algorithms. Primarily based on a C++ API, it supports real-time image and video processing across various platforms.

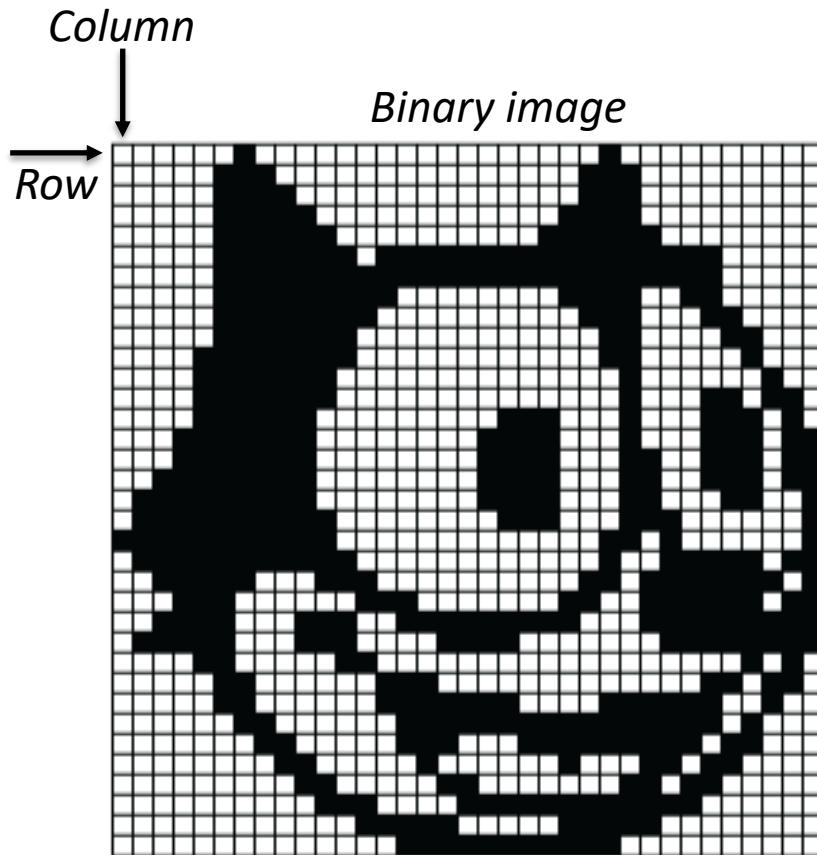
Key Modules:

- **core** – Basic data structures (e.g., Mat) and functions.
- **imgproc** – Image processing (filters, transformations, color conversion).
- **video** – Motion analysis, background subtraction, tracking.
- **calib3d** – Camera calibration, pose estimation, 3D reconstruction.
- **features2d** – Feature detection, description, and matching.
- **objdetect** – Object detection (e.g., faces, people, cars).
- **highgui** – Simple GUI for image/video display.
- **videoio** – Video capture and codec interface.
- **Others** – Python bindings, FLANN, testing tools, etc.

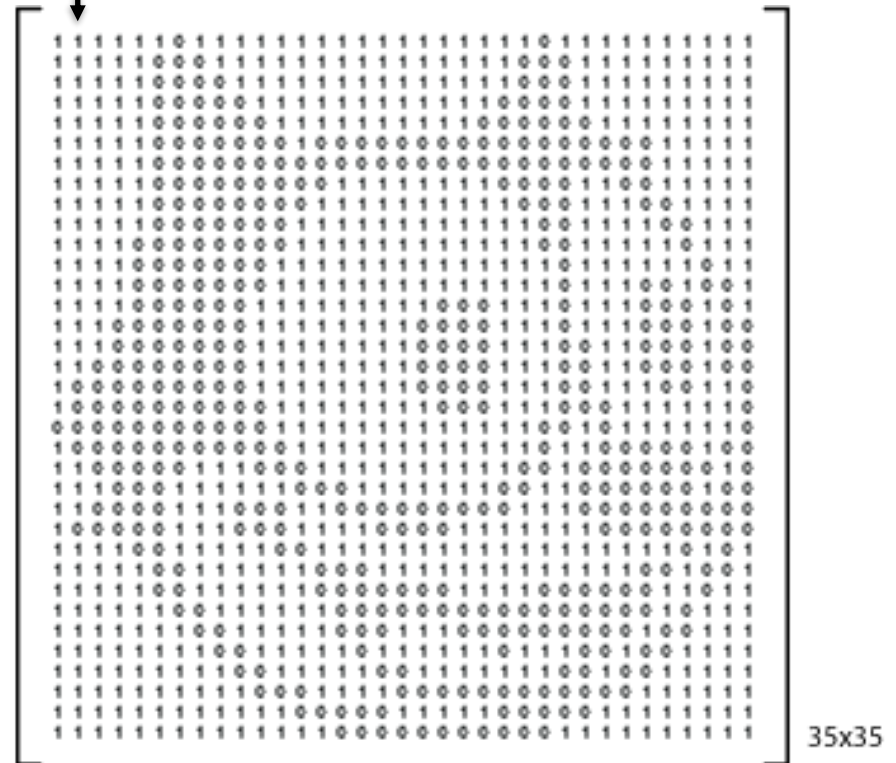
More info: opencv.org

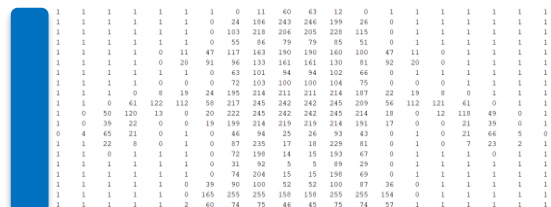
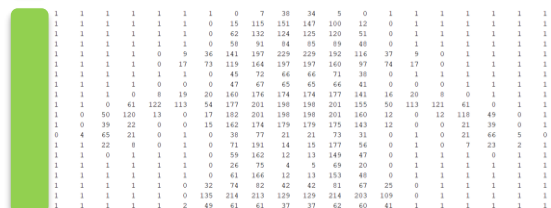
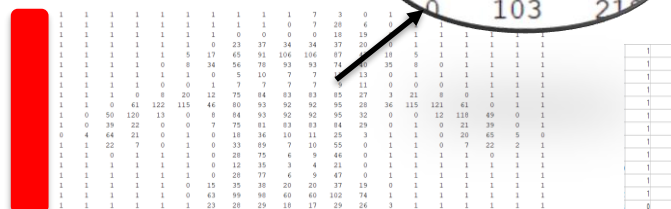
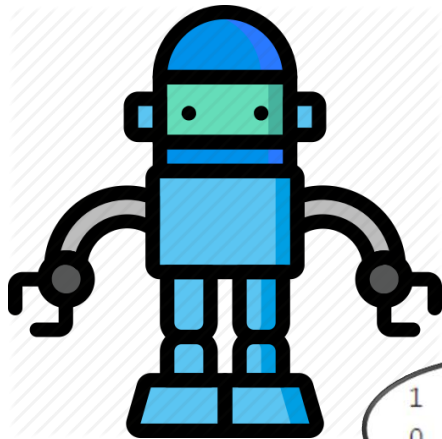


2D image matrix

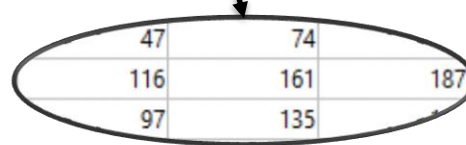
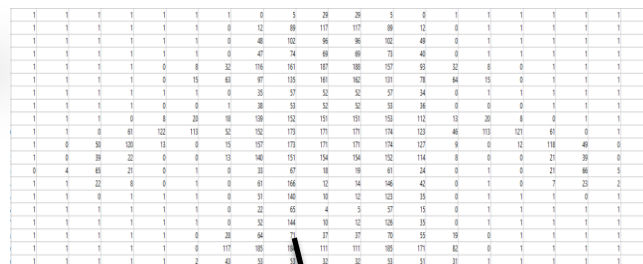
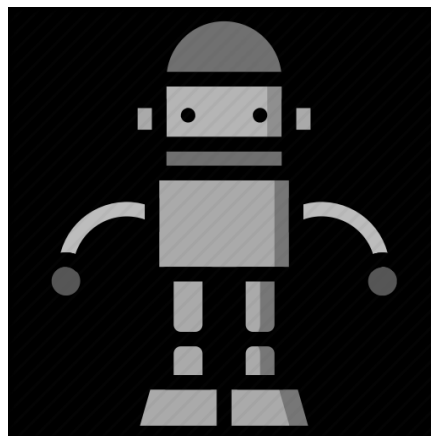


1bit (0-1)

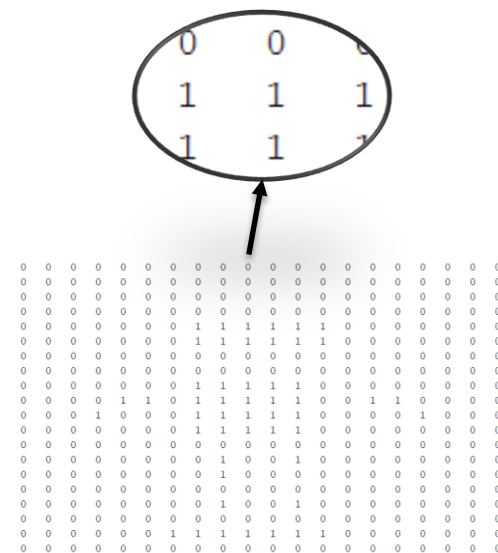
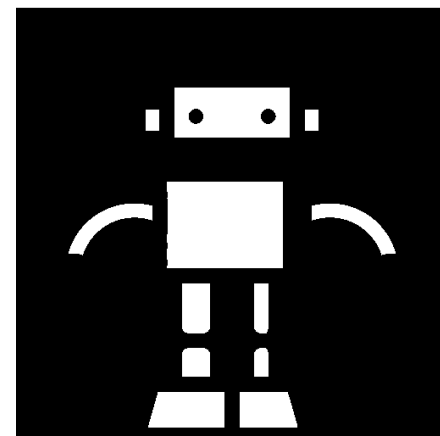




Color 3 channels



Grayscale 1 channel



Monochrome 1 channel

OpenCV minimum project example

CmakeLists.txt

```
cmake_minimum_required(VERSION 3.5)
project(OpenCV)

set(CMAKE_CXX_STANDARD 17)
set(CMAKE_CXX_STANDARD_REQUIRED ON)
set(CMAKE_CXX_EXTENSIONS OFF)

find_package( OpenCV REQUIRED )
include_directories( ${OpenCV_INCLUDE_DIRS} )

add_executable(OpenCV src/test.cpp)

target_link_libraries(OpenCV ${OpenCV_LIBS} )
```



test.cpp

```
#include <opencv2/opencv.hpp>
#include <iostream>
int main() {
    // Print OpenCV version
    std::cout << "OpenCV version: " << CV_VERSION <<
    std::endl;
    // Create a simple image
    cv::Mat image = cv::Mat::zeros(300, 300, CV_8UC3);
    // Draw a red circle
    cv::circle(image, cv::Point(150, 150), 50, cv::Scalar(0, 0,
    255), -1);
    // Show the image
    cv::imshow("OpenCV Test", image);
    cv::waitKey(0); // Wait for key press

    return 0;
}
```

Basic variables and commands in OpenCV

Mat → Image matrix variable

- `Imread(„image_name”,1) → read from file`
- `Declare matrix of size and type → Mat::zeros(Size(500, 500), CV_8U)`

`Mat A, C; // creates just the header parts`

`A = imread(“image_name”, IMREAD_COLOR); // here we'll know the method used (allocate matrix)`

`Mat B(A); // Use the copy constructor`

`C = A; // Assignment operator`

All the objects, in the end, point to the same single data matrix and making a modification using any of them will affect all the other ones as well.

`Mat F = A.clone();`

`Mat G;`

`A.copyTo(G);`

Now modifying F or G will not affect the matrix pointed by the A's header

`Imshow(„window_name”,mat_name) → show image matrix in a window`

`waitKey()` → **The function waitKey waits for a key event infinitely (when `delay ≤ 0`) or for delay milliseconds, when it is positive**

Color conversion → `cvtColor(image,image,COLOR_RGB2GRAY);`

`resize()` → Resizes an image.

`resize (InputArray src, OutputArray dst, Size dsize, double fx = 0, double fy = 0, int interpolation = INTER_LINEAR)`

Use OpenCV documentation → <https://docs.opencv.org/4.5.5/d1/dfb/intro.html>

OpenCV documentation - examples

Mat → https://docs.opencv.org/4.5.5/d3/d63/classcv_1_1Mat.html#details

Detailed Description

n-dimensional dense array class

The class `Mat` represents an n-dimensional dense numerical single-channel or multi-channel array. It can be used to store real or complex-valued vectors and matrices, grayscale or color images, voxel volumes, vector fields, point clouds, tensors, histograms (though, very high-dimensional histograms may be better stored in a `SparseMat`). The data layout of the array `M` is defined by the array `M.step[]`, so that the address of element $(i_0, \dots, i_{M_{\text{dims}}-1})$, where $0 \leq i_k < M.\text{size}[k]$, is computed as:

$$\text{addr}(M_{i_0, \dots, i_{M_{\text{dims}}-1}}) = M.\text{data} + M.\text{step}[0] * i_0 + M.\text{step}[1] * i_1 + \dots + M.\text{step}[M.\text{dims} - 1] * i_{M_{\text{dims}}-1}$$

In case of a 2-dimensional array, the above formula is reduced to:

$$\text{addr}(M_{i,j}) = M.\text{data} + M.\text{step}[0] * i + M.\text{step}[1] * j$$

Note that `M.step[i] >= M.step[i+1]` (in fact, `M.step[i] >= M.step[i+1]*M.size[i+1]`). This means that 2-dimensional matrices are stored row-by-row, 3-dimensional matrices are stored plane-by-plane, and so on. `M.step[M.dims-1]` is minimal and always equal to the element size `M.elemSize()`.

So, the data layout in `Mat` is compatible with the majority of dense array types from the standard toolkits and SDKs, such as Numpy (ndarray), Win32 (independent device bitmaps), and others, that is, with any array that uses `steps` (or `strides`) to compute the position of a pixel. Due to this compatibility, it is possible to make a `Mat` header for user-allocated data and process it in-place using OpenCV functions.

There are many different ways to create a `Mat` object. The most popular options are listed below:

`imshow(„window_name”,mat_name) →`

https://docs.opencv.org/4.5.5/d7/dfc/group__highgui.html#ga453d42fe4cb60e5723281a89973ee563

◆ imshow()

```
void cv::imshow ( const String & winname,
                  InputArray  mat
                  )
```

Python:

```
cv.imshow( winname, mat ) -> None
```

```
#include <opencv2/highgui.hpp>
```

Displays an image in the specified window.

The function `imshow` displays an image in the specified window. If the window was created with the `cv::WINDOW_AUTOSIZE` flag, the image is shown with its original size, however it is still limited by the screen resolution. Otherwise, the image is scaled to fit the window. The function may scale the image, depending on its depth:

- If the image is 8-bit unsigned, it is displayed as is.
- If the image is 16-bit unsigned, the pixels are divided by 256. That is, the value range `[0,255*256]` is mapped to `[0,255]`.
- If the image is 32-bit or 64-bit floating-point, the pixel values are multiplied by 255. That is, the value range `[0,1]` is mapped to `[0,255]`.
- 32-bit integer images are not processed anymore due to ambiguity of required transform. Convert to 8-bit unsigned matrix using a custom preprocessing specific to image's context.

If window was created with OpenGL support, `cv::imshow` also support `ogl::Buffer`, `ogl::Texture2D` and `cuda::GpuMat` as input.

If the window was not created before this function, it is assumed creating a window with `cv::WINDOW_AUTOSIZE`.

If you need to show an image that is bigger than the screen resolution, you will need to call `namedWindow("", WINDOW_NORMAL)` before the `imshow`.

Use OpenCV documentation → <https://docs.opencv.org/4.5.5/d1/dfb/intro.html>



University of
Zagreb



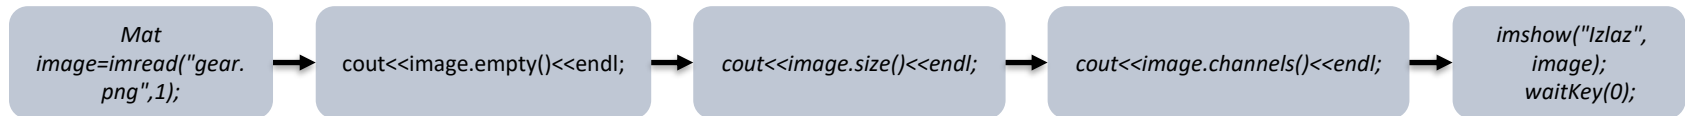
Faculty of mechanical
engineering and naval
architecture

Practical task 1

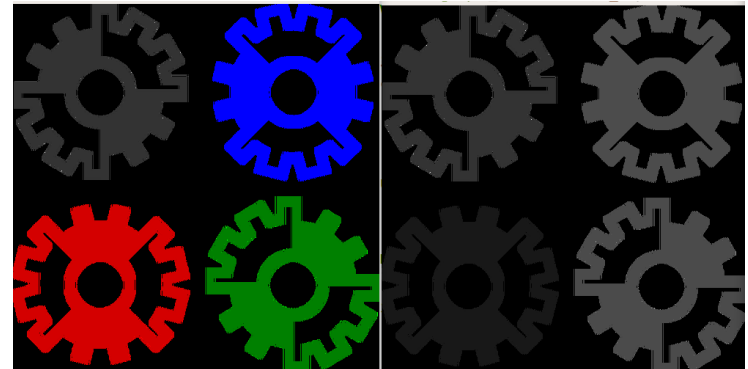
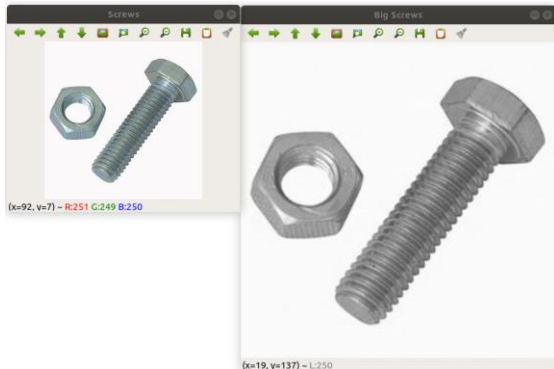
1. Create black image of size (10,10). Write out the content of Mat, size and number of channels. Show the image in the output window.



2. Read images from files (gear.png and screws.jpg) and show them in the window. Write out empty(), size and number of channels for the image read from the file.



3. Copy (clone) mat of gear.png, convert color to grayscale, resize (twice the size) and show image.

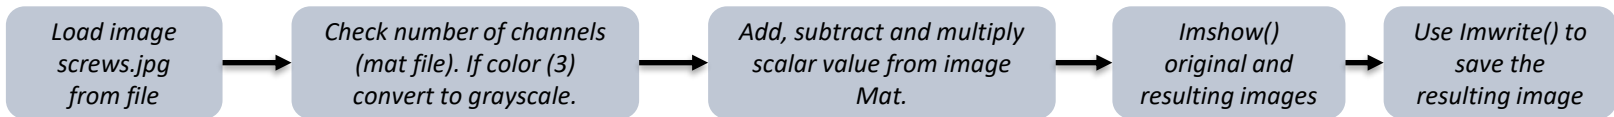


Practical task 2

Basic processing on grayscale images:

- Brighten, darken and contrast grayscale image.

Version 1



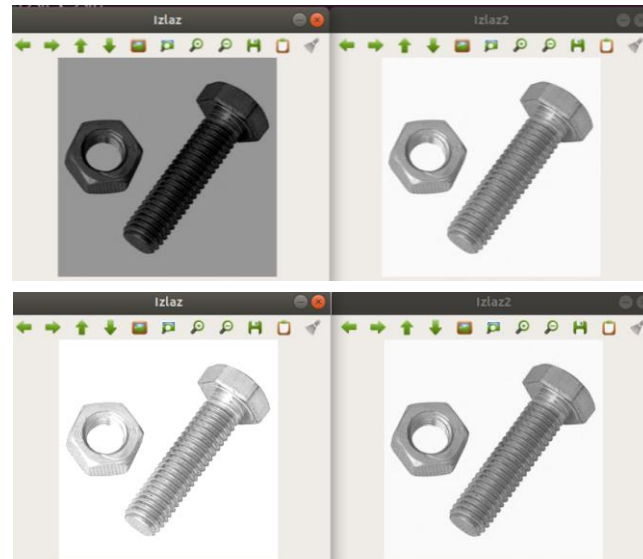
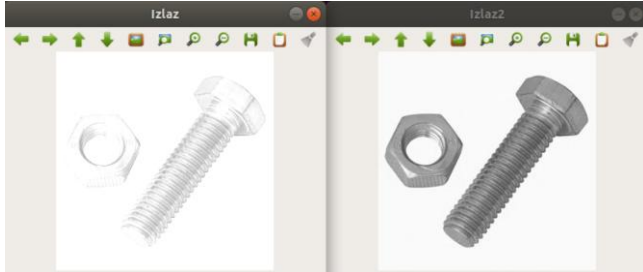
Comment – error if values out of scope of matrix!!!

Version 2



`add(image,Mat::ones(image.size(), CV_8U)*100,image);`

`subtract(image,Mat::ones(image.size(), CV_8U)*100,image);`



`image.convertTo(image,-1,0.5,0);`

$$g(i,j) = \alpha \cdot f(i,j) + \beta$$

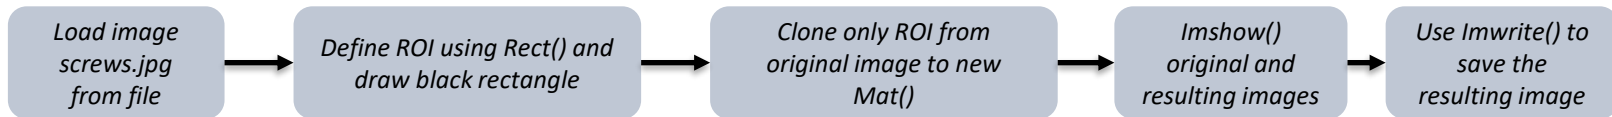
File name

Variable of type Mat()

`imwrite("processed_image.jpg",image);`

Practical task 3

Extract rectangular ROI from the image:



// Setup a rectangle to define your region of interest

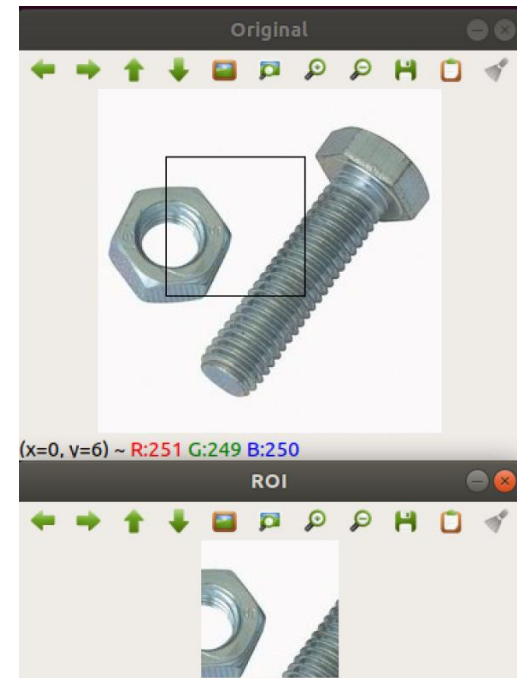
```
Rect myROI(50, 50, 100, 100);
```

Upper left row Width Height

Upper left column

```
Mat image_ROI=image(myROI).clone();
```

Crop and clone the original image with previously defined ROI



Practical task 4

Draw circle, rectangle and line:

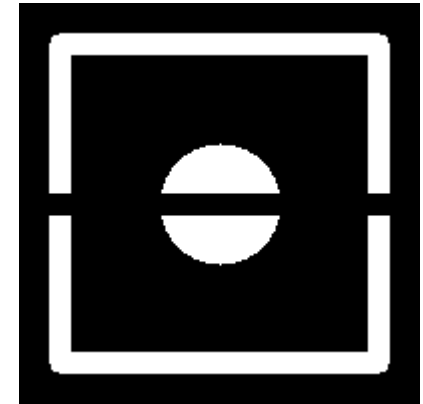
Version 1

Construct grayscale
black image,
Size(200,200)

Draw white circle(full),
white rectangle rectangle
(not full) and black line

Show image

```
Mat image = Mat::zeros(Size(200, 200), CV_8U);  
circle(image, Point(100,100),30,Scalar(255),-1,LINE_8,0);  
rectangle(image, Rect(20,20,160,160),Scalar(255),10,LINE_8,0);  
line(image,Point(0,100),Point(200,100),Scalar(0),10,LINE_8,0);
```



```
+ circle()
void cv::circle ( InputOutputArray img,
                  Point center,
                  int radius,
                  const Scalar & color,
                  int thickness = 1,
                  int lineType = LINE_8,
                  int shift = 0
                )
Python:
cv.circle( img, center, radius, color[, thickness[, lineType[, shift]]] ) -> img

#include <opencv2/imgproc.hpp>

Draws a circle.
The function cv::circle draws a simple or filled circle with a given center and radius.

Parameters
img      Image where the circle is drawn.
center   Center of the circle.
radius   Radius of the circle.
color    Circle color.
thickness Thickness of the circle outline, if positive. Negative values, like FILLED, mean that a filled circle is to be drawn.
lineType Type of the circle boundary. See LineTypes
shift    Number of fractional bits in the coordinates of the center and in the radius value.

Examples:
samples/cpp/converthull.cpp, samples/cpp/floodcolor.cpp, samples/cpp/kmeans.cpp, samples/cpp/ldemo.cpp,
samples/cpp/minarea.cpp, samples/cpp/tutorial_code/imgProc/basic_drawing/Drawimg_1.cpp,
samples/cpp/tutorial_code/imgProc/basic_drawing/Drawimg_2.cpp, samples/cpp/tutorial_code/imgTrans/houghcircles.cpp, and
samples/dnn/openpose.cpp.
```

```
+ rectangle() [1/2]
void cv::rectangle ( InputOutputArray img,
                    Point pt1,
                    Point pt2,
                    const Scalar & color,
                    int thickness = 1,
                    int lineType = LINE_8,
                    int shift = 0
                  )
Python:
cv.rectangle( img, pt1, pt2, color[, thickness[, lineType[, shift]]] ) -> img
cv.rectangle( img, rec, color[, thickness[, lineType[, shift]]] ) -> img

#include <opencv2/imgproc.hpp>

Draws a simple, thick, or filled up-right rectangle.
The function cv::rectangle draws a rectangle outline or a filled rectangle whose two opposite corners are pt1 and pt2.

Parameters
img      Image.
pt1      Vertex of the rectangle.
pt2      Vertex of the rectangle opposite to pt1.
color    Rectangle color or brightness (grayscale image).
thickness Thickness of lines that make up the rectangle. Negative values, like FILLED, mean that the function has to draw a filled rectangle.
lineType Type of the line. See LineTypes
shift    Number of fractional bits in the point coordinates.
```

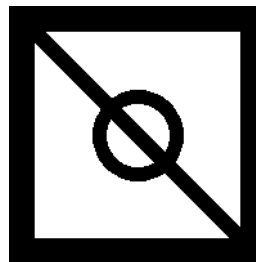
```
+ line()
void cv::line ( InputOutputArray img,
               Point pt1,
               Point pt2,
               const Scalar & color,
               int thickness = 1,
               int lineType = LINE_8,
               int shift = 0
             )
Python:
cv.line( img, pt1, pt2, color[, thickness[, lineType[, shift]]] ) -> img

#include <opencv2/imgproc.hpp>

Draws a line segment connecting two points.
The function line draws the line segment between pt1 and pt2 points in the image. The line is clipped by the image boundaries. For non-antialiased lines with integer coordinates, the 8-connected or 4-connected Bresenham algorithm is used. Thick lines are drawn with rounding endings. Antialiased lines are drawn using Gaussian filtering.

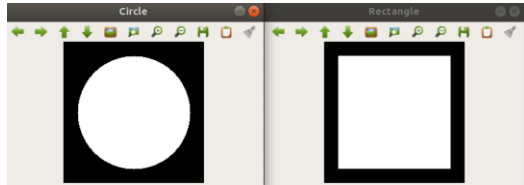
Parameters
img      Image.
pt1      First point of the line segment.
pt2      Second point of the line segment.
color    Line color.
thickness Line thickness.
lineType Type of the line. See LineTypes.
shift    Number of fractional bits in the point coordinates.
```

Student exercise. Do it yourself →



Practical task 5

Test bitwise operands on binary circle and rectangle:



```
void circle(InputOutputArray img, Point center, int radius, const Scalar &color, int thickness = 1, int lineType = LINE_8, int shift = 0)
circle(circle1,cv::Point(100,100),80,Scalar(255),-1,8,0)

void rectangle(InputOutputArray img, Rect rec, const Scalar &color, int thickness = 1, int lineType = LINE_8, int shift = 0)
rectangle(rectangle1,Rect(20,20,160,160),Scalar(255),-1,8,0);
```

Construct 2 grayscale
black images,
Size(200,200)

Draw white circle and
rectangle in separate Mats

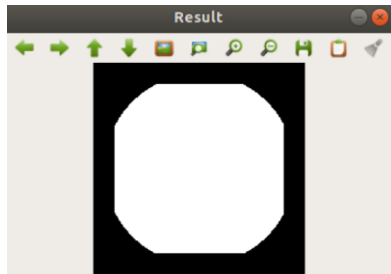
Test bitwise_and(),
bitwise_or(),
bitwise_xor(),...

Show images

Concatenate and
save resulting
images

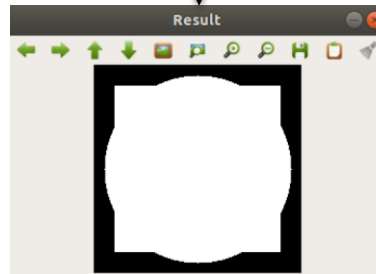
Bitwise AND

```
void bitwise_and(InputArray src1, InputArray src2, OutputArray dst, InputArray mask = noArray())
bitwise_and(circle1,rectangle1,result);
```



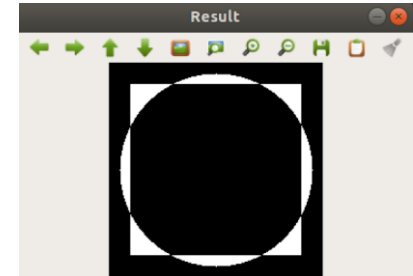
```
void bitwise_or(InputArray src1, InputArray src2, OutputArray dst, InputArray mask = noArray())
bitwise_or(circle1,rectangle1,result);
```

Bitwise OR



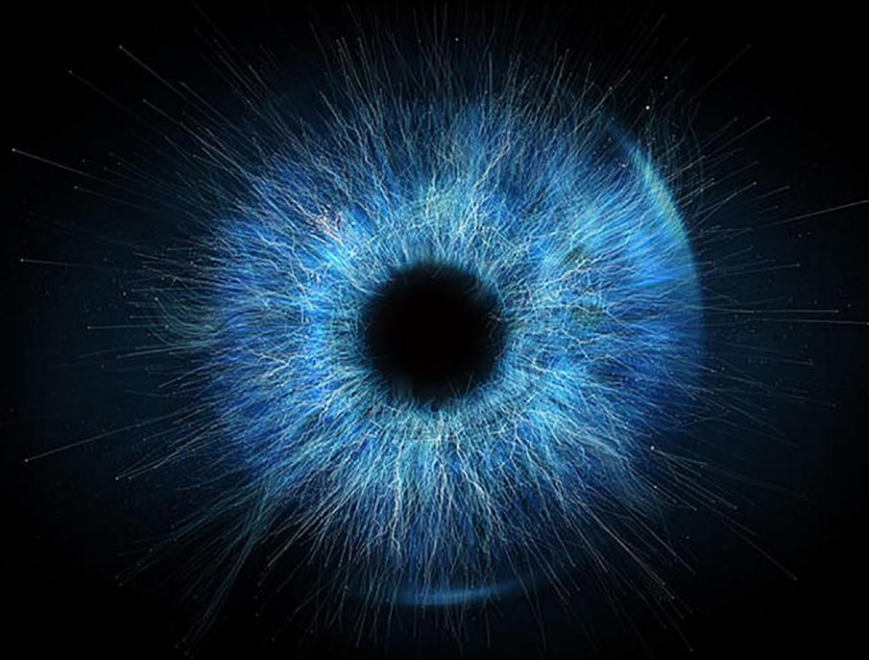
Bitwise XOR

```
void bitwise_xor(InputArray src1, InputArray src2, OutputArray dst, InputArray mask = noArray())
bitwise_xor(circle1,rectangle1,result);
```



```
void hconcat(InputArrayOfArrays src, OutputArray dst)
hconcat(image,image,connected);
```

Horizontal concatenating of
image matrices



Exercise – Filters

Machine vision algorithms

Doc.dr.sc. Filip Šuligoj
fsuligoj@fsb.hr



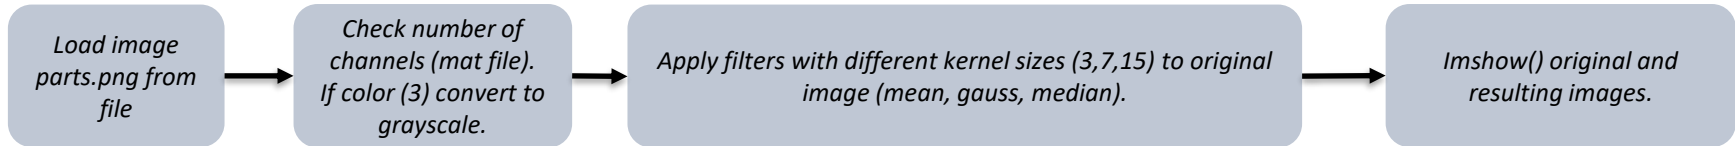
University of
Zagreb



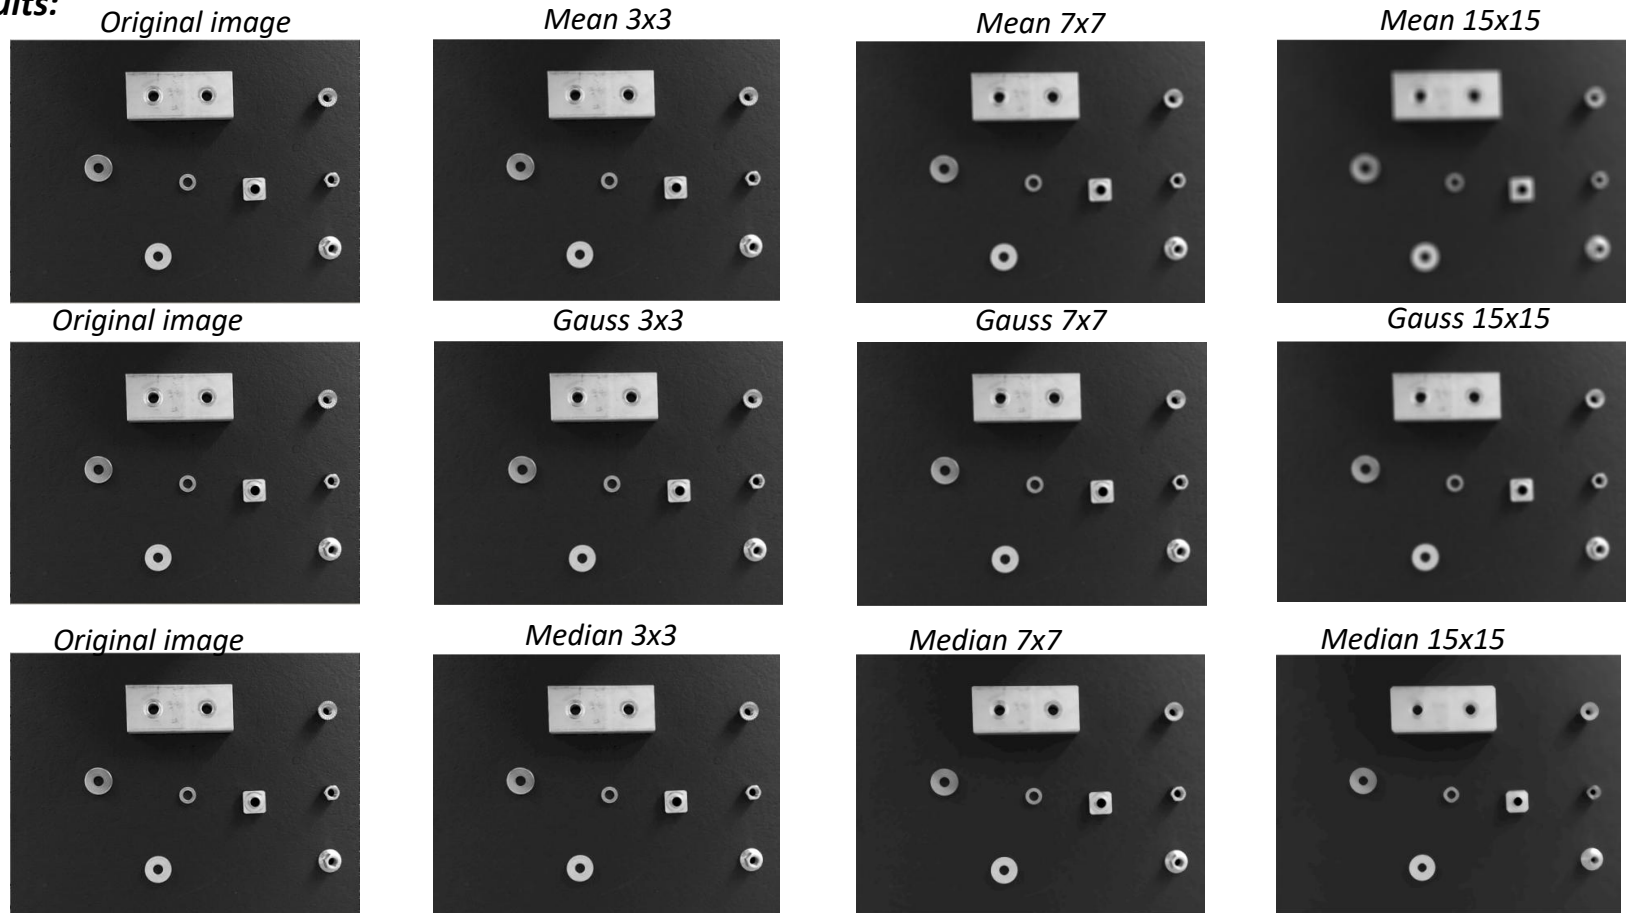
Faculty of mechanical
engineering and naval
architecture

Practical task 1

Apply image filters to the original images



Results:



Practical task 1

https://docs.opencv.org/master/d4/d86/group__imgproc__filter.html#ga8c45db9afe636703801b0b2e440fce37

◆ blur()

```
void cv::blur ( InputArray  src,
                OutputArray dst,
                Size        ksize,
                Point        anchor = Point(-1,-1) ,
                int          borderType = BORDER_DEFAULT
              )
```

Python:

```
dst = cv.blur( src, ksize[, dst[, anchor[, borderType]]])
```

```
#include <opencv2/imgproc.hpp>
```

Blurs an image using the normalized box filter.

The function smooths an image using the kernel:

$$K = \frac{1}{\text{ksize.width} * \text{ksize.height}} \begin{bmatrix} 1 & 1 & 1 & \dots & 1 & 1 \\ 1 & 1 & 1 & \dots & 1 & 1 \\ \dots & & & & & \\ 1 & 1 & 1 & \dots & 1 & 1 \end{bmatrix}$$

The call `blur(src, dst, ksize, anchor, borderType)` is equivalent to `boxFilter(src, dst, src.type(), ksize, anchor, true, borderType)`.

Parameters

- src** input image; it can have any number of channels, which are processed independently, but the depth should be CV_8U, CV_16U, CV_16S, CV_32F or CV_64F.
- dst** output image of the same size and type as src.
- ksize** blurring kernel size.
- anchor** anchor point; default value Point(-1,-1) means that the anchor is at the kernel center.
- borderType** border mode used to extrapolate pixels outside of the image, see [BorderTypes](#). **BORDER_WRAP** is not supported.

See also

[boxFilter](#), [bilateralFilter](#), [GaussianBlur](#), [medianBlur](#)

Examples:

[samples/cpp/edge.cpp](#), [samples/cpp/laplace.cpp](#), and [samples/cpp/tutorial_code/ImgProc/Smoothing/Smoothing.cpp](#).



Practical task 1

https://docs.opencv.org/master/d4/d86/group__imgproc__filter.html#gaabe8c836e97159a9193fb0b11ac52cf1

◆ GaussianBlur()

```
void cv::GaussianBlur ( InputArray  src,
                       OutputArray dst,
                       Size        ksize,
                       double       sigmaX,
                       double       sigmaY = 0 ,
                       int          borderType = BORDER_DEFAULT
                     )
```

Python:

```
dst = cv.GaussianBlur( src, ksize, sigmaX[, dst[, sigmaY[, borderType]]] )
```

```
#include <opencv2/imgproc.hpp>
```

Blurs an image using a Gaussian filter.

The function convolves the source image with the specified Gaussian kernel. In-place filtering is supported.

Parameters

- src** input image; the image can have any number of channels, which are processed independently, but the depth should be CV_8U, CV_16U, CV_16S, CV_32F or CV_64F.
- dst** output image of the same size and type as src.
- ksize** Gaussian kernel size. ksize.width and ksize.height can differ but they both must be positive and odd. Or, they can be zero's and then they are computed from sigma.
- sigmaX** Gaussian kernel standard deviation in X direction.
- sigmaY** Gaussian kernel standard deviation in Y direction; if sigmaY is zero, it is set to be equal to sigmaX, if both sigmas are zeros, they are computed from ksize.width and ksize.height, respectively (see [getGaussianKernel](#) for details); to fully control the result regardless of possible future modifications of all this semantics, it is recommended to specify all of ksize, sigmaX, and sigmaY.
- borderType** pixel extrapolation method, see [BorderTypes](#). **BORDER_WRAP** is not supported.

See also

[sepFilter2D](#), [filter2D](#), [blur](#), [boxFilter](#), [bilateralFilter](#), [medianBlur](#)

Examples:

[samples/cpp/laplace.cpp](#), [samples/cpp/tutorial_code/imgProc/Smoothing/Smoothing.cpp](#), and [samples/cpp/tutorial_code/imgTrans/Sobel_Demo.cpp](#).



Practical task 1

https://docs.opencv.org/master/d4/d86/group__imgproc__filter.html#ga564869aa33e58769b4469101aac458f9

◆ medianBlur()

```
void cv::medianBlur ( InputArray  src,
                     OutputArray dst,
                     int         ksize
                   )
```

Python:

```
dst = cv.medianBlur( src, ksize[, dst] )
```

```
#include <opencv2/imgproc.hpp>
```

Blurs an image using the median filter.

The function smoothes an image using the median filter with the **ksize** × **ksize** aperture. Each channel of a multi-channel image is processed independently. In-place operation is supported.

Note

The median filter uses **BORDER_REPLICATE** internally to cope with border pixels, see **BorderTypes**

Parameters

src input 1-, 3-, or 4-channel image; when ksize is 3 or 5, the image depth should be CV_8U, CV_16U, or CV_32F, for larger aperture sizes, it can only be CV_8U.

dst destination array of the same size and type as src.

ksize aperture linear size; it must be odd and greater than 1, for example: 3, 5, 7 ...

See also

[bilateralFilter](#), [blur](#), [boxFilter](#), [GaussianBlur](#)

Examples:

[samples/cpp/laplace.cpp](#), [samples/cpp/tutorial_code/imgProc/Smoothing/Smoothing.cpp](#), and [samples/cpp/tutorial_code/imgTrans/houghcircles.cpp](#).



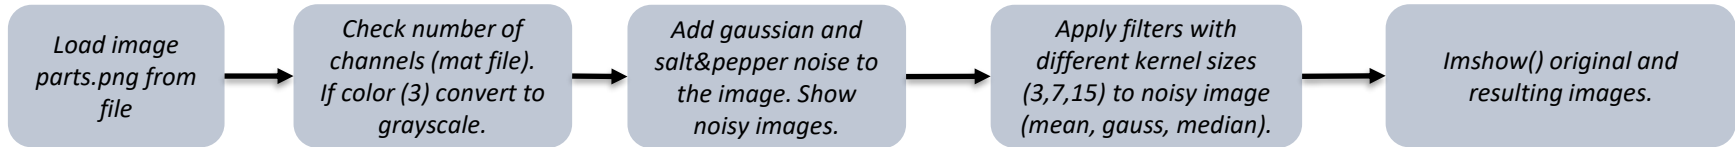
University of
Zagreb



Faculty of mechanical
engineering and naval
architecture

Practical task 2

Add noise (gaussian and salt&pepper) to image



Adding gaussian noise to the image:

```
Mat gaussian_noise = Mat(image.size(),image.type());
randn(gaussian_noise,50,10);
add(image,gaussian_noise,image);
```

void cv::randn (InputOutputArray dst, InputArray mean, InputArray stddev)
Fills the array with normally distributed random numbers. The function cv::randn fills the matrix dst with normally distributed random numbers with the specified mean vector and the standard deviation matrix. The generated random numbers are clipped to fit the value range of the output array data type.

void cv::randu (InputOutputArray dst, InputArray low, InputArray high)
Generates a single uniformly-distributed random number or an array of random numbers. Non-template variant of the function fills the matrix dst with uniformly-distributed random numbers from the specified range.

Adding salt&pepper noise to the image:

```
Mat saltpepper_noise = Mat::zeros(image.rows, image.cols,CV_8U);
randu(saltpepper_noise,0,255);
```

```
Mat black = saltpepper_noise < 30;
Mat white = saltpepper_noise > 225;
```

```
image.setTo(255,white);
image.setTo(0,black);
```

`yy.setTo(0)` will set all the pixels to 0.

`yy.setTo(0, xx)` will set all the pixels who have a corresponding pixel with a non-zero value in the `xx` `Mat` to 0.

Example:

```
yy =
2 2 2
2 2 2
2 2 2
```

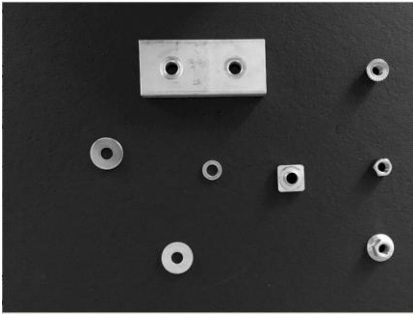
```
xx =
0 0 0
0 1 0
0 0 0
```

`yy.setTo(0, xx) =>`

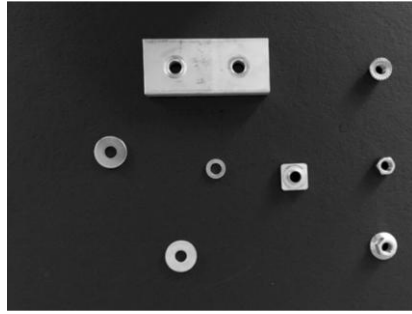
```
yy =
2 2 2
2 0 2
2 2 2
```

Practical task 2 - results

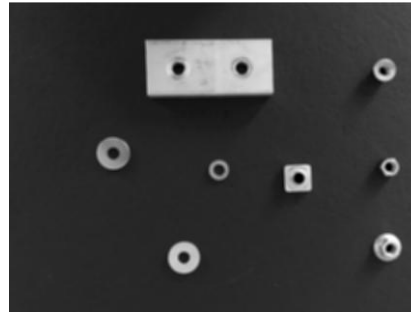
Original image



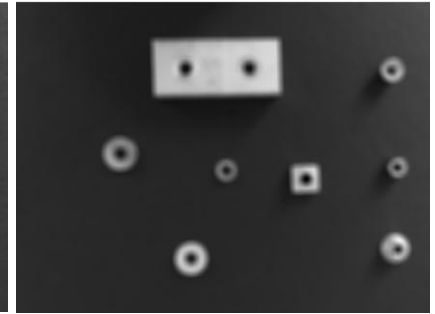
Mean 3x3



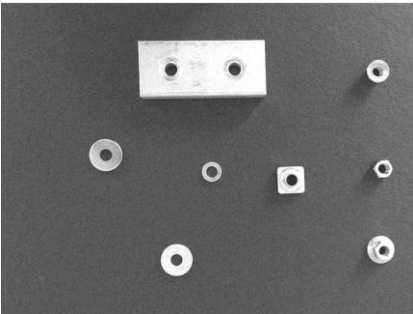
Mean 7x7



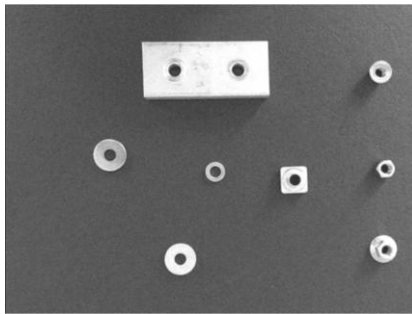
Mean 15x15



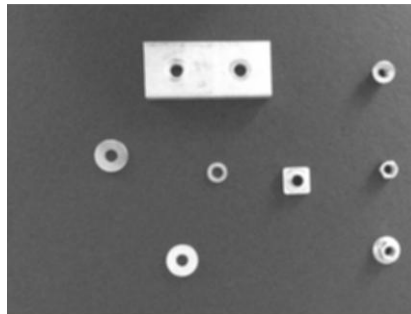
Gaussian noise



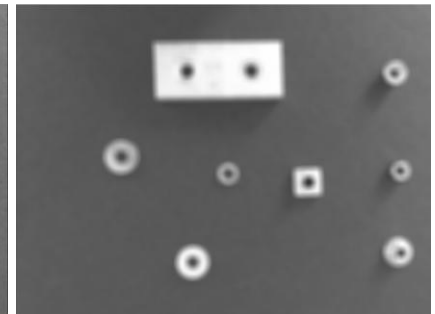
Mean 3x3



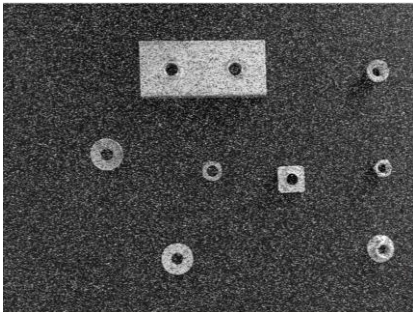
Mean 7x7



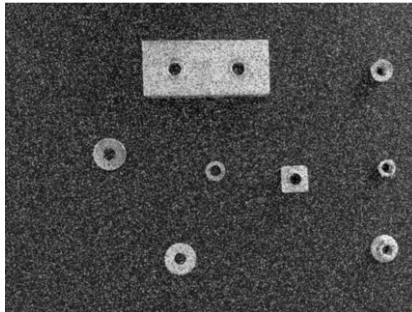
Mean 15x15



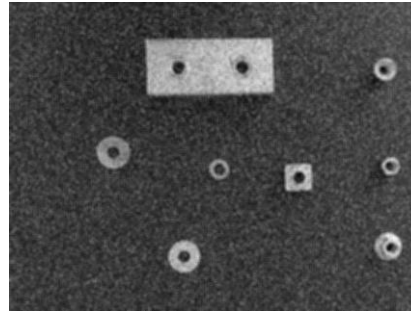
Salt and pepper noise



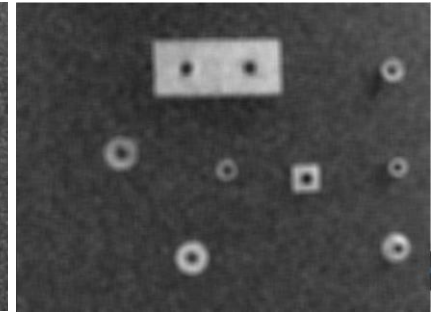
Mean 3x3



Mean 7x7

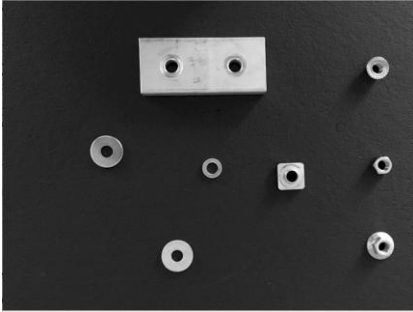


Mean 15x15

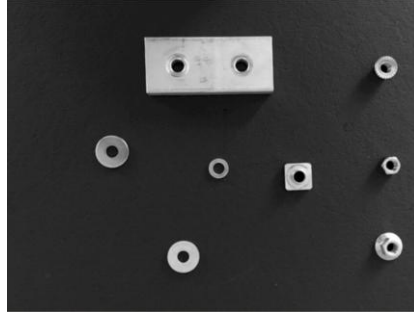


Practical task 2 - results

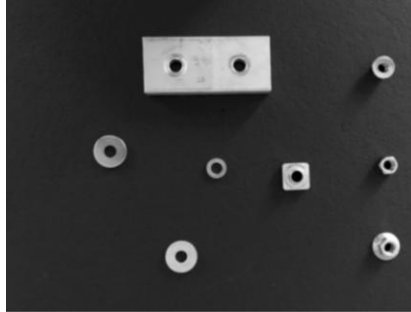
Original image



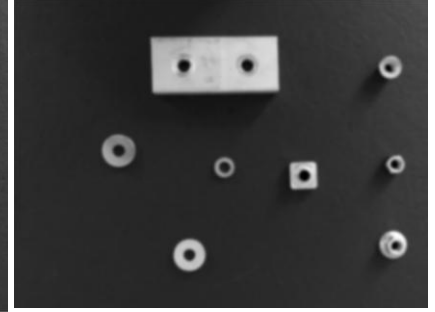
Gauss 3x3



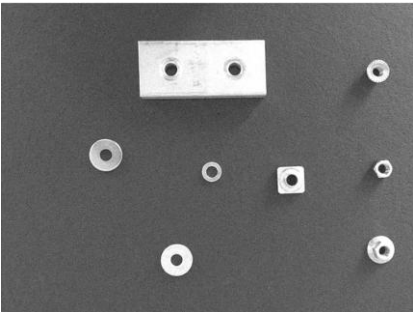
Gauss 7x7



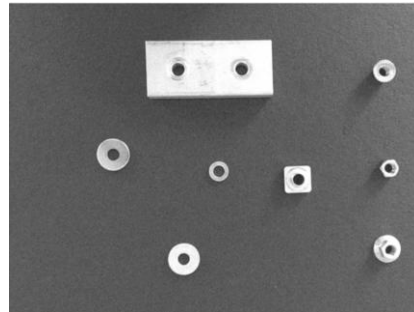
Gauss 15x15



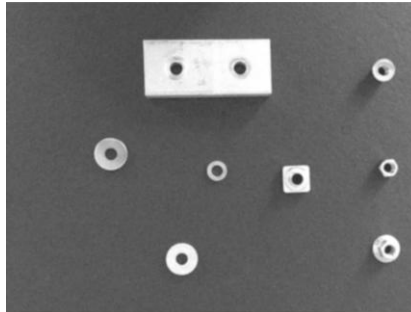
Gaussian noise



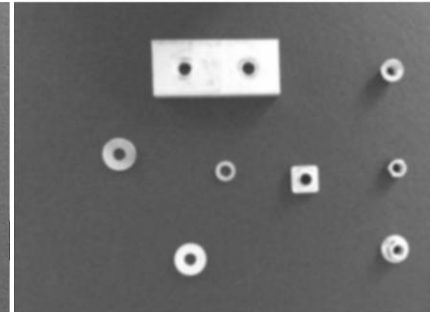
Gauss 3x3



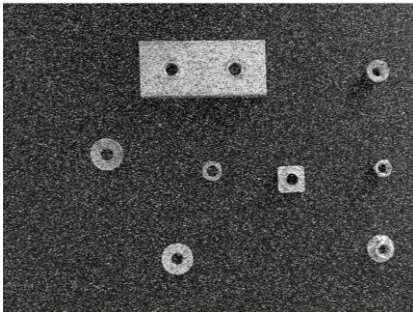
Gauss 7x7



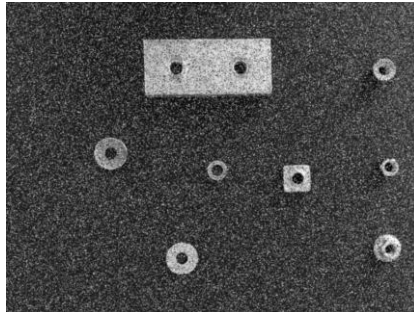
Gauss 15x15



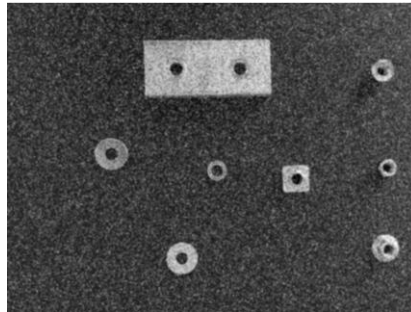
Salt and pepper noise



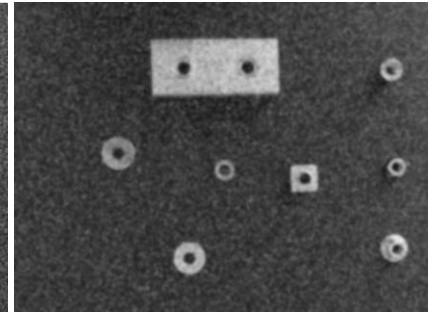
Gauss 3x3



Gauss 7x7

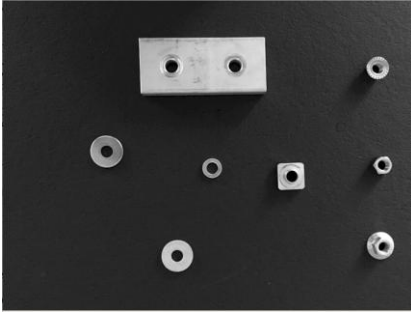


Gauss 15x15

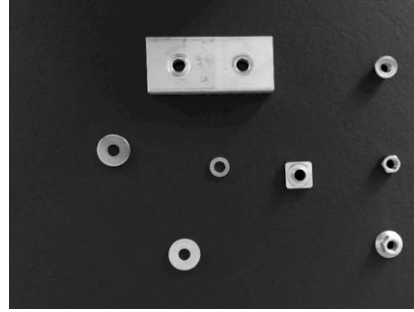


Practical task 2 - results

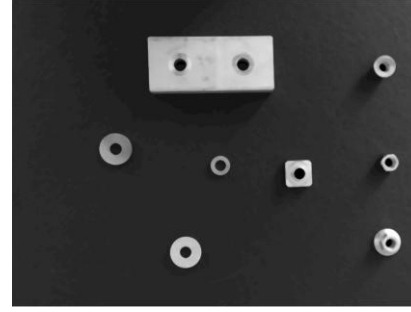
Original image



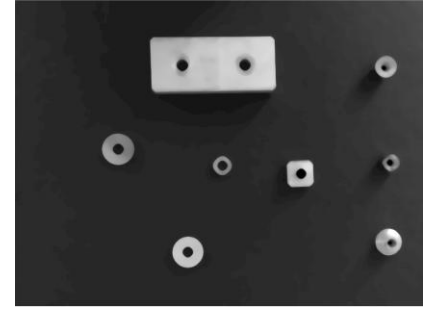
Median 3x3



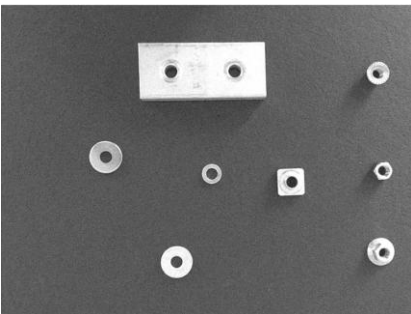
Median 7x7



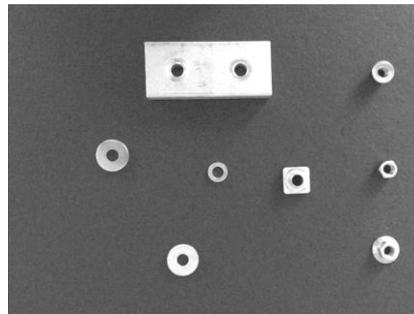
Median 15x15



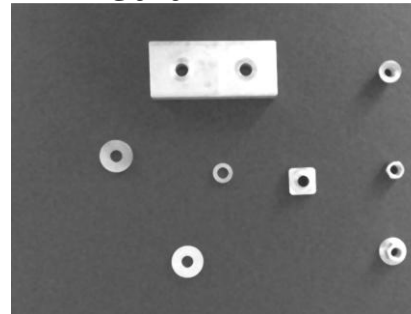
Gaussian noise



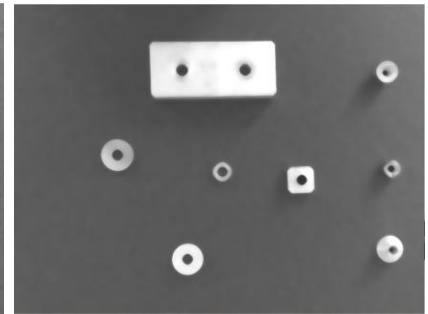
Median 3x3



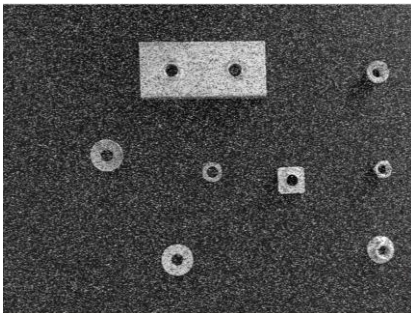
Median 7x7



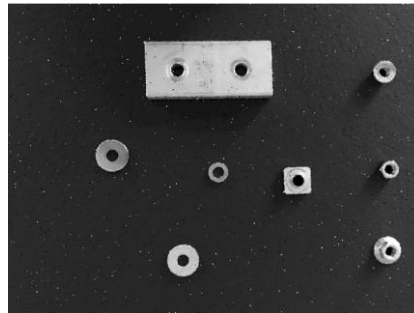
Median 15x15



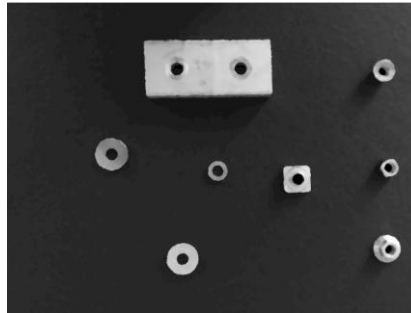
Salt and pepper noise



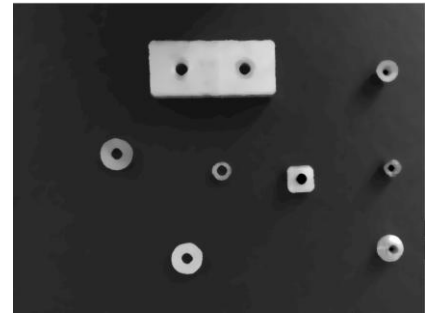
Median 3x3

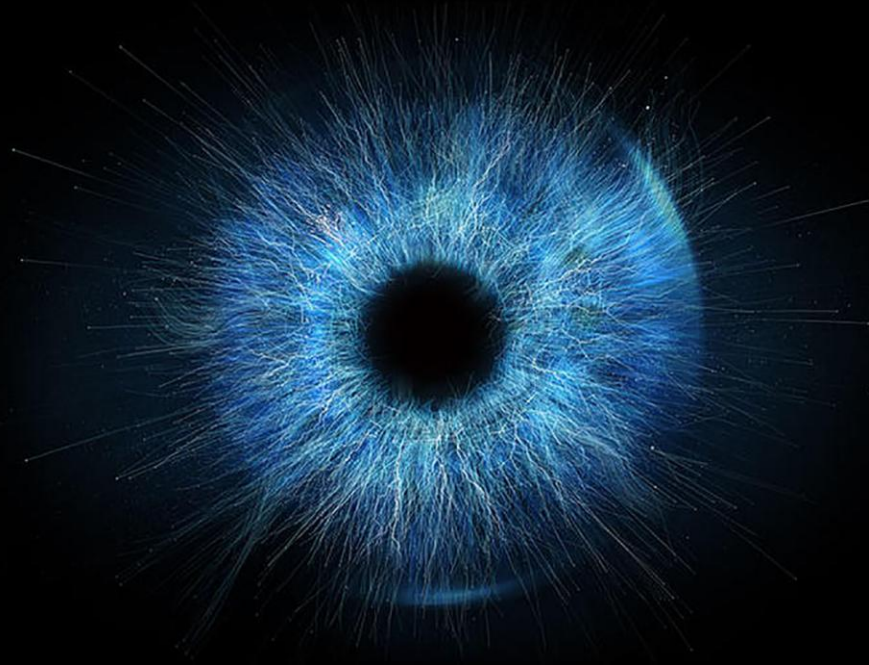


Median 7x7



Median 15x15





Exercise – Image morphology

Fundamentals of machine vision algorithms

dr.sc. Filip Šuligoj
fsuligoj@fsb.hr



University of
Zagreb

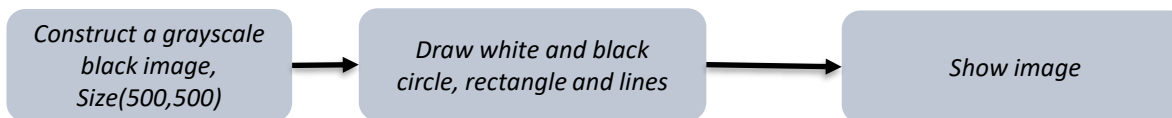


Faculty of mechanical
engineering and naval
architecture

Practical task 1

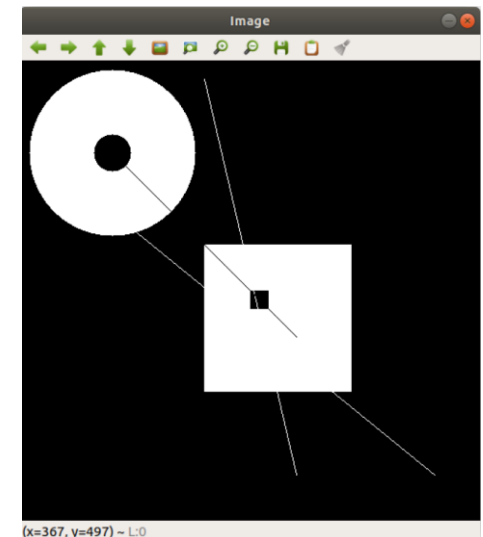
Test morphological operators → erosion, dilatation, opening, closing, gradient, tophat, blackhat on the image with intersecting shapes.

Step 1 (create grayscale image with intersecting shapes)



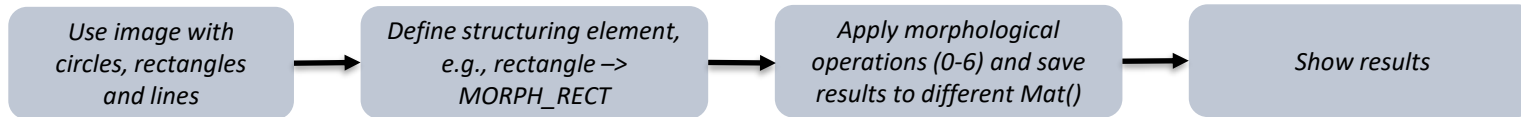
```
circle(image,cv::Point(100,100),90,Scalar(255),-1,8,0);
rectangle(image,Rect(200,200,160,160),Scalar(255),-1,8,0);
line(image,cv::Point(200,20),cv::Point(300,450),Scalar(255),1,8,0);
```

Starting point Ending point Color/Intensity Line type
Line thickness



Practical task 1

Test morphological operators → erosion, dilatation, opening, closing, gradient, tophat, blackhat on the image with intersecting shapes:



```
int element_type=MORPH_RECT; → Structuring element type
int element_size=1; → Structuring element size
Mat element = getStructuringElement( element_type,
                                   Size( 2*element_size+1, 2*element_size+1 ),
                                   Point( element_size, element_size ) );
```

```
//erode( image, result, element);
//dilate(image,result2,element);
morphologyEx(image,result,0,element);
morphologyEx(image,result2,1,element);
morphologyEx(image,result3,2,element);
morphologyEx(image,result4,3,element);
morphologyEx(image,result5,4,element);
morphologyEx(image,result6,5,element);
morphologyEx(image,result7,6,element);
```

morphologyEx types →

0	MORPH_ERODE	see erode
1	MORPH_DILATE	see dilate
2	MORPH_OPEN	an opening operation
		dst=open(src,element)=dilate(erode(src,element))
3	MORPH_CLOSE	a closing operation
		dst=close(src,element)=erode(dilate(src,element))
4	MORPH_GRADIENT	a morphological gradient
		dst=morph_grad(src,element)=dilate(src,element)-erode(src,element)
5	MORPH_TOPHAT	"top hat"
		dst=tophat(src,element)=src-open(src,element)
6	MORPH_BLACKHAT	"black hat"
		dst=blackhat(src,element)=close(src,element)-src
7	MORPH_HITMISS	"hit or miss" - Only supported for CV_8UC1 binary images. A tutorial can be found in the documentation

◆ morphologyEx()

```
void cv::morphologyEx( InputArray src,
                      OutputArray dst,
                      int op,
                      InputArray kernel,
                      Point anchor = Point(-1,-1),
                      int iterations = 1,
                      int borderType = BORDER_CONSTANT,
                      const Scalar & borderValue = morphologyDefaultBorderValue()
                      )
```

Python:

```
cv.morphologyEx( src, op, kernel[, dst[, anchor[, iterations[, borderType[, borderValue]]]]) -> dst
```

The function `cv::morphologyEx` can perform advanced morphological transformations using an erosion and dilation as basic operations.

Any of the operations can be done in-place. In case of multi-channel images, each channel is processed independently.

Parameters

- src** Source image. The number of channels can be arbitrary. The depth should be one of CV_8U, CV_16U, CV_16S, CV_32F or CV_64F.
- dst** Destination image of the same size and type as source image.
- op** Type of a morphological operation, see [MorphTypes](#)
- kernel** Structuring element. It can be created using [getStructuringElement](#).
- anchor** Anchor position with the kernel. Negative values mean that the anchor is at the kernel center.
- iterations** Number of times erosion and dilation are applied.
- borderType** Pixel extrapolation method, see [BorderTypes](#). **BORDER_WRAP** is not supported.
- borderValue** Border value in case of a constant border. The default value has a special meaning.



University of
Zagreb



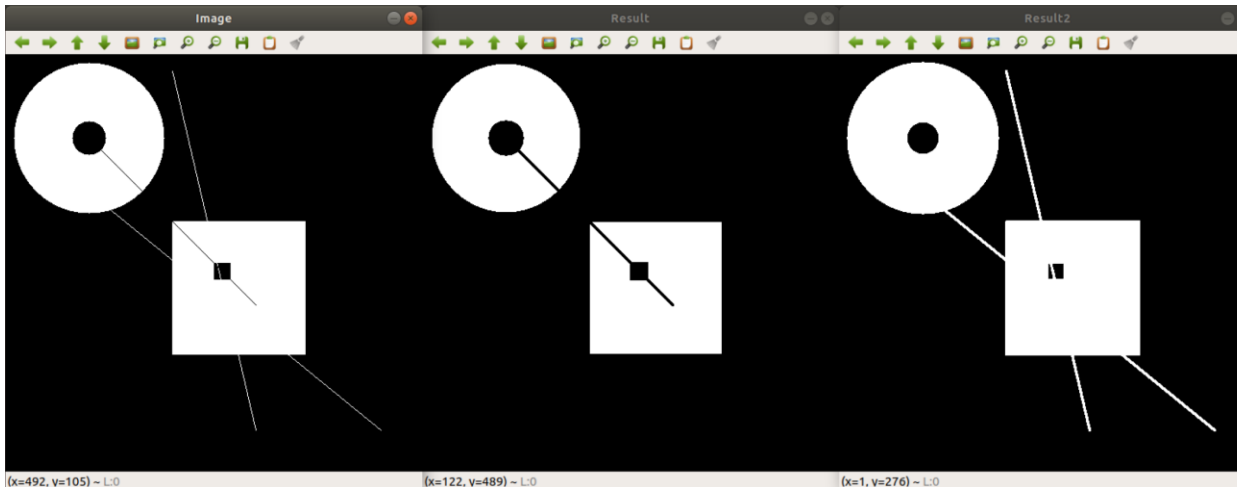
Faculty of mechanical
engineering and naval
architecture

Practical task 1 - results

Original

erosion

dilatation



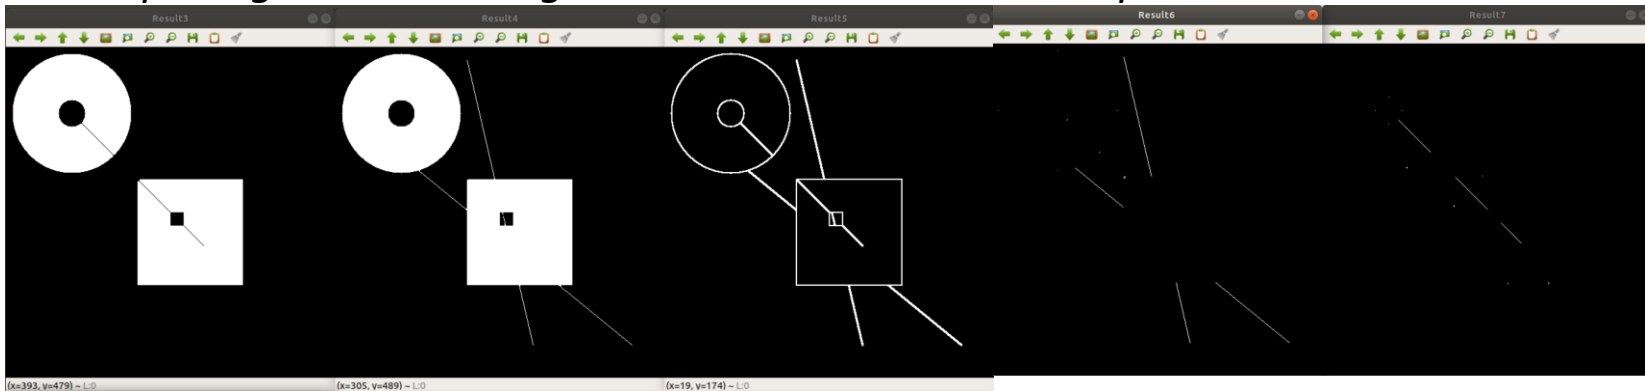
opening

closing

gradient

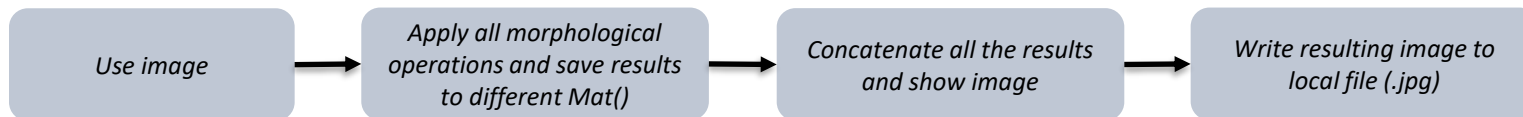
tophat

blackhat



Practical task 2

Test erosion, dilatation, opening, closing, gradient, tophat, blackhat with different structuring elements:



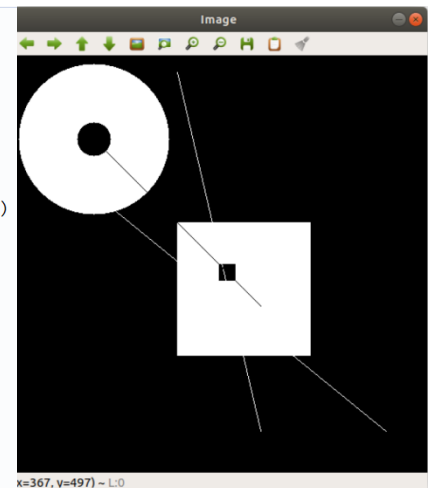
Test parameters:

- Different structuring elements sizes (3x3, 5x5, 7x7)
- Different structuring element types (rect, ellipse, cross)
- Different line thicknesses (1, 2, 3, 4 pixel wide) in the original image.

```
# Rectangular Kernel
>>> cv.getStructuringElement(cv.MORPH_RECT,(5,5))
array([[1, 1, 1, 1, 1],
       [1, 1, 1, 1, 1],
       [1, 1, 1, 1, 1],
       [1, 1, 1, 1, 1],
       [1, 1, 1, 1, 1]], dtype=uint8)

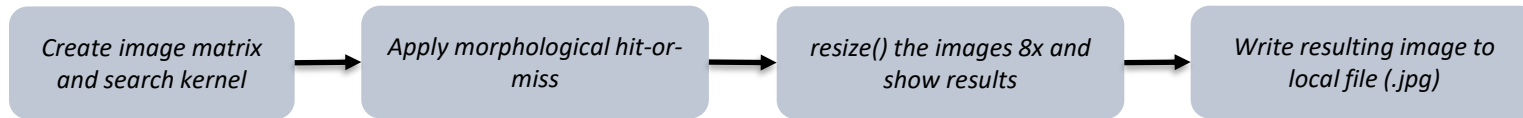
# Elliptical Kernel
>>> cv.getStructuringElement(cv.MORPH_ELLIPSE,(5,5))
array([[0, 0, 1, 0, 0],
       [1, 1, 1, 1, 1],
       [1, 1, 1, 1, 1],
       [1, 1, 1, 1, 1],
       [0, 0, 1, 0, 0]], dtype=uint8)

# Cross-shaped Kernel
>>> cv.getStructuringElement(cv.MORPH_CROSS,(5,5))
array([[0, 0, 1, 0, 0],
       [0, 0, 1, 0, 0],
       [1, 1, 1, 1, 1],
       [0, 0, 1, 0, 0],
       [0, 0, 1, 0, 0]], dtype=uint8)
```



Practical task 3

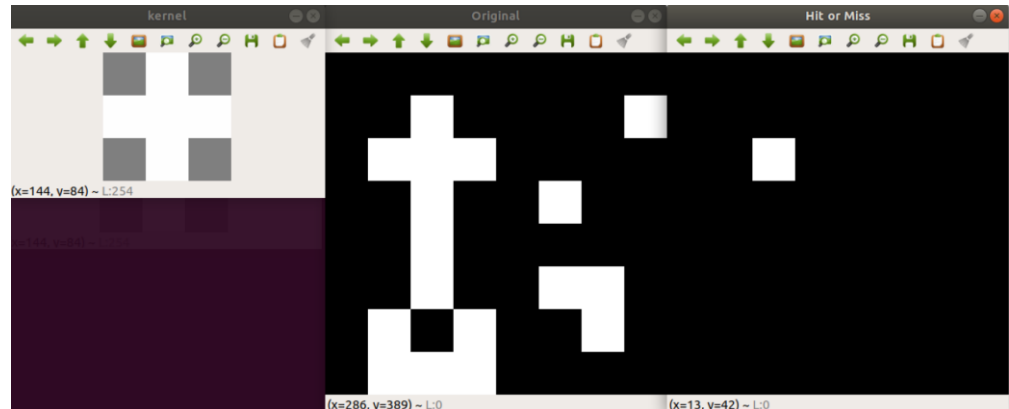
Use morphological hit-or-miss to find a shape of a cross:



```
morphologyEx(input_image, output_image, MORPH_HITMISS, kernel);
```

```
Mat input_image = (Mat_<uchar>(8, 8) <<
    0, 0, 0, 0, 0, 0, 0, 0,
    0, 0, 255, 0, 0, 0, 0, 255,
    0, 255, 255, 255, 0, 0, 0, 0,
    0, 0, 255, 0, 0, 255, 0, 0,
    0, 0, 255, 0, 0, 0, 0, 0,
    0, 0, 255, 0, 0, 255, 255, 0,
    0, 255, 0, 255, 0, 0, 255, 0,
    0, 255, 255, 255, 0, 0, 0, 0);
Mat kernel = (Mat_<int>(3, 3) <<
    0, 1, 0,
    1, 1, 1,
    0, 1, 0);
```

Result



Student assignment (seminar)



University of
Zagreb



Faculty of mechanical
engineering and naval
architecture

Student assignment

Add trackbar and change the filter kernel size manually in code from Practical task 2. Add and test out the bilateral filter as well.

For trackbar use example :

https://docs.opencv.org/master/da/d6a/tutorial_trackbar.html

◆ bilateralFilter()

```
void cv::bilateralFilter ( InputArray src,
                          OutputArray dst,
                          int d,
                          double sigmaColor,
                          double sigmaSpace,
                          int borderType = BORDER_DEFAULT
                        )
```

Python:

```
dst = cv.bilateralFilter( src, d, sigmaColor, sigmaSpace[, dst[, borderType]] )
```

```
#include <opencv2/imgproc.hpp>
```

Applies the bilateral filter to an image.

The function applies bilateral filtering to the input image, as described in http://www.dai.ed.ac.uk/CVonline/LOCAL_COPIES/MANDUCHI1/Bilateral_Filtering.html bilateralFilter can reduce unwanted noise very well while keeping edges fairly sharp. However, it is very slow compared to most filters.

Sigma values: For simplicity, you can set the 2 sigma values to be the same. If they are small (< 10), the filter will not have much effect, whereas if they are large (> 150), they will have a very strong effect, making the image look "cartoonish".

Filter size: Large filters ($d > 5$) are very slow, so it is recommended to use $d=5$ for real-time applications, and perhaps $d=9$ for offline applications that need heavy noise filtering.

This filter does not work inplace.

Parameters

src Source 8-bit or floating-point, 1-channel or 3-channel image.

dst Destination image of the same size and type as src .

d Diameter of each pixel neighborhood that is used during filtering. If it is non-positive, it is computed from sigmaSpace.

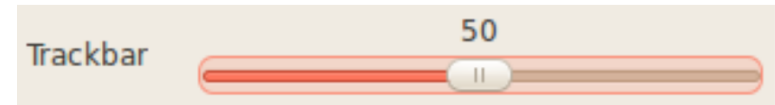
sigmaColor Filter sigma in the color space. A larger value of the parameter means that farther colors within the pixel neighborhood (see sigmaSpace) will be mixed together, resulting in larger areas of semi-equal color.

sigmaSpace Filter sigma in the coordinate space. A larger value of the parameter means that farther pixels will influence each other as long as their colors are close enough (see sigmaColor). When $d > 0$, it specifies the neighborhood size regardless of sigmaSpace. Otherwise, d is proportional to sigmaSpace.

borderType border mode used to extrapolate pixels outside of the image, see [BorderTypes](#)

Examples:

```
samples/cpp/tutorial_code/imgProc/Smoothing/Smoothing.cpp.
```



https://docs.opencv.org/master/d4/d86/group_imgproc_filter.html#ga9d7064d478c95d60003cf839430737ed



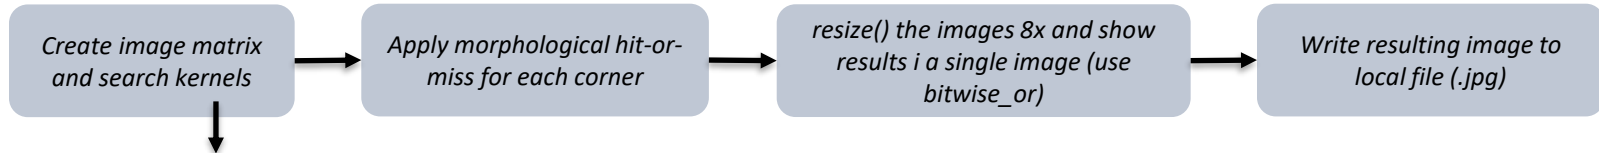
University of
Zagreb



Faculty of mechanical
engineering and naval
architecture

Student assignment

Use morphological hit-or-miss to find 4 corners of a square or a rectangle:



Hint: search kernels should also have negative fields!

Example:

